

# SAMRO Artist Radio-Activity Dashboard pre-processing and testing

## Import relevant packages for preprocessing

In [147]:

```
import numpy as np
import pandas as pd
import scipy.stats as stats
from scipy.stats import shapiro
import matplotlib.pyplot as plt
import pylab
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.stattools import adfuller, kpss

from adtk.detector import SeasonalAD
from adtk.data import validate_series

import warnings
warnings.filterwarnings('ignore')

%matplotlib inline
```

## import CSV file and output data

In [11]:

```
data_original = pd.read_excel("Black Coffee - 2020.xlsx")
data = data_original.copy()
data.head()
```

Out[11]:

	Station	Artist	Title	Imp's	Sec	Date/time
0	North West FM	Black Coffee feat. Nakhane Toure	We Dance Again	21,000	0	01 Jan 2020 - 00:07
1	MFM	Black Coffee feat. Shekhinah	Your Eyes	0	0	01 Jan 2020 - 00:08
2	947	Black Coffee & David Guetta feat. Delilah Montagu	Drive	11,000	180	01 Jan 2020 - 00:35
3	Eldos FM	Black Coffee feat. Soulstar	You Rock My World	0	0	01 Jan 2020 - 00:39
4	YFM 99.2	Black Coffee feat. Zhao	Any Other Way	29,000	0	01 Jan 2020 - 00:50

## Pre-processing of data

In [12]:

```

# Remove leading white spce
for x in data.columns:
    data.rename(columns={x:x.strip()}, inplace=True)

#rename columns to make it easier to read
data.rename(columns={"Imp's": 'Impacts', 'Date/time': 'Date'}, inplace=True)

# remove all channels where impacts = 0 as they are assumed not to have official RAM
data = data[data.Impacts != 0]

#convert date column to a date object to enable Time Series analysis
data['Date'] = pd.to_datetime(data['Date'])

#Sort columns by date
data.sort_values(by='Date', inplace=True)

#Convert Impacts into numeric format
data['Impacts'] = data['Impacts'].str.replace(',', '')
data['Impacts'] = data['Impacts'].str.extract('(\d+)', expand=False) # remove all non-digits
data['Impacts'] = data['Impacts'].map(lambda x: pd.to_numeric(x)) #convert to integer

# Fill null values and inf values with 0
data['Impacts'].fillna(0, inplace=True)
data.replace([np.inf, -np.inf], 0, inplace=True)

#convert 'Seconds' and 'Impacts' columns to integer types
data["Sec"] = pd.to_numeric(data["Sec"])
data['Impacts'] = data['Impacts'].astype('int64')

#Get rid of rows that dont make sense ie where there are no second of play but have
#this could be because of incorrect data capture or under 1 second of play is recorded
data = data[~((data["Impacts"] > 0) & (data["Sec"] == 0))]

```

## Exploratory code

In [13]:

```

#count number of unique Radio Stations tracked
data.Station.nunique()

```

Out[13]:

43

In [14]:

```
# Check overview of dataset
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 8970 entries, 2 to 22716
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Station    8970 non-null   object
 1   Artist     8970 non-null   object
 2   Title      8970 non-null   object
 3   Impacts    8970 non-null   int64
 4   Sec        8970 non-null   int64
 5   Date       8970 non-null   datetime64[ns]
dtypes: datetime64[ns](1), int64(2), object(3)
memory usage: 490.5+ KB
```

In [15]:

```
#Get basic summary statistics of dataset
data.describe()
```

Out[15]:

	Impacts	Sec
<b>count</b>	8970.000000	8970.000000
<b>mean</b>	138217.505463	249.258194
<b>std</b>	157437.301387	93.966489
<b>min</b>	0.000000	0.000000
<b>25%</b>	20000.000000	180.000000
<b>50%</b>	79333.000000	240.000000
<b>75%</b>	203000.000000	312.000000
<b>max</b>	998000.000000	972.000000

In [16]:

```
# Check date ranges of data, since it is one full year we will not need to subset it
print (data.Date.min())
print (data.Date.max())
```

```
2020-01-01 00:35:00
2020-12-31 23:08:00
```

In [17]:

```
original_data = data.copy() #make a copy of the pre-processed data as a backup  
data.tail()
```

Out[17]:

	Station	Artist	Title	Impacts	Sec	Date
22705	Munghana Lonene FM	Black Coffee feat. Sabrina Claudio	Sbcnscsly	584000	432	2020-12-31 19:44:00
22706	Ligwalagwala FM	Black Coffee feat. Sabrina Claudio	Sbcnscsly	348000	354	2020-12-31 20:16:00
22708	Jacaranda	Black Coffee feat. Usher	LaLaLa	233000	204	2020-12-31 20:31:00
22714	KFM	Black Coffee feat. Nakhane Toure	We Dance Again	212000	324	2020-12-31 22:09:00
22716	SAfm	Black Coffee feat. Nakhane Toure	We Dance Again	13000	330	2020-12-31 23:08:00

## Check for a normal distribution

In [18]:

```
# Check for normality of dataframe for 'Impacts'


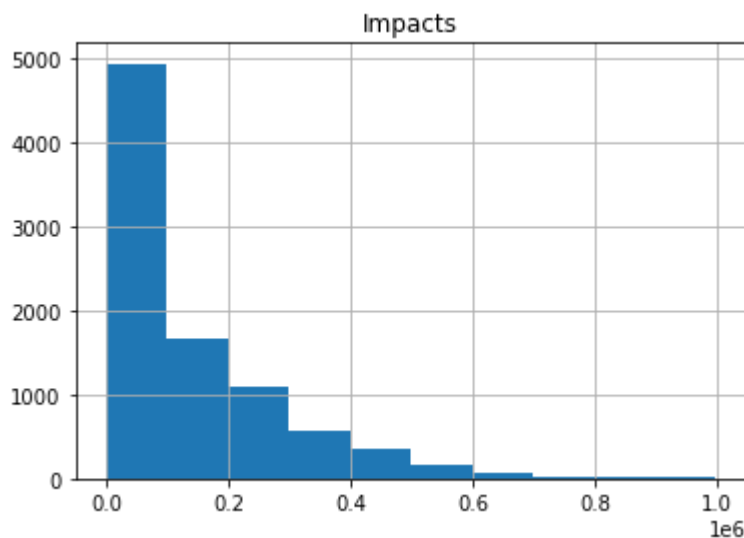
# Mean () and Median () not close together
dt = data[['Impacts']]
print(f'Kurtosis: {dt.kurt()}')
print(f'Skewness: {dt.skew()}')

# Do a Shapiro Wilkes test
stat, p = shapiro(dt.iloc[:, :].values)
print(f'Shapiro-Wilkes: Statistic: {stat}, p-value = {p}')

data.hist('Impacts')
plt.show()

#Create a QQ plot to check for normal distribution
stats.probplot(data['Impacts'], dist="norm", plot=pylab)
pylab.show()
```

```
Kurtosis: Impacts      3.768057
dtype: float64
Skewness: Impacts      1.763089
dtype: float64
Shapiro-Wilkes: Statistic: 0.808011531829834, p-value = 0.0
```



In [19]:

```

# Check for normality of dataframe for 'seconds'
# Mean (249) and Median (240) are close together
# Data looks like it follows a normal distribution
d = data[['Sec']]
print(f'Kurtosis: {d.kurt()}')
print(f'Skewness: {d.skew()}')

# Do a Shapiro Wilkes test
statd, pd = shapiro(d)
print(f'Shapiro-Wilkes: Statistic: {statd}, p-value = {pd}')

data.hist('Sec')
plt.show()



#Create a QQ plot to check for normal distribution
stats.probplot(data['Sec'], dist="norm", plot=pylab)
pylab.show()

```

```

Kurtosis: Sec      4.473664
dtype: float64
Skewness: Sec      1.057147
dtype: float64
Shapiro-Wilkes: Statistic: 0.935672402381897, p-value = 0.0

```



It seems that 'Impacts' does not follow a normal distribution and is skewed right. 'Secs' is very slightly skewed

right but resembles a normal distribution. The Quantile Quantile plot (QQ) of the 'Secs' variable show that it is close to a normal distribution. I will assume as such.

## Use Bartlett's test to test for change in variance before and after lockdown and album launch

In [20]:

```
# Split data into two datasets, to check seconds played data before and after the release date
# Release date: 1 May 2020

PA = data[['Sec', 'Date']]
PreAlbum = PA[PA['Date'] < '2020-05-01']

PostAlbum = PA[PA['Date'] >= '2020-05-01']

#Remove Date component to check for variance of Seconds played values within timeframe

PreAlbumSec = PreAlbum['Sec']
PostAlbumSec = PostAlbum['Sec']
```

In [21]:

```
PreAlbumSec.describe()
```

Out[21]:

```
count      2118.000000
mean        224.059490
std         87.013311
min          0.000000
25%         180.000000
50%         198.000000
75%         264.000000
max         930.000000
Name: Sec, dtype: float64
```

In [22]:

```
PostAlbumSec.describe()
```

Out[22]:

```
count      6852.000000
mean        257.047285
std         94.672360
min          0.000000
25%         180.000000
50%         252.000000
75%         330.000000
max         972.000000
Name: Sec, dtype: float64
```

In [23]:

```
# Do Bartlett's test to check significance of variance between two groups
stats.bartlett(PreAlbumSec, PostAlbumSec)
```

Out[23]:

```
BartlettResult(statistic=22.328772996158293, pvalue=2.297384464247441e-06)
```

Since p value is less than 0.05 there is evidence that the the change in variance is significant after the release of his album release in 2020

In [24]:

```
#bartlett's test for equal variance before and after the first lockdown date that started

PreLockdown = PA[PA['Date'] < '2020-03-26']

PostLockdown = PA[PA['Date'] >= '2020-03-26']
PostAlbum.head()

#Remove Date component to check for variance of Seconds played values within timeframe

PreLockdownSec = PreLockdown['Sec']
PostLockdownSec = PostLockdown['Sec']
# Do Bartlett's test to check significance of variance between two groups
stats.bartlett(PreLockdownSec, PostLockdownSec)
```

Out[24]:

```
BartlettResult(statistic=66.36816061913828, pvalue=3.740962870588646e-16)
```

Since p value is less than 0.05 there is evidence that the the change in variance is significant after the lockdown started.

## T-Test for change in means after lockdown

Do a t-test to test for significance of means before and after the lockdown started on Black Coffee's seconds of play

$H_0$ : Lockdown had no effect on Black Coffee's spins after lockdown was announced

$H_1$ : Lockdown had an effect on Black Coffee's spins after lockdown was announced



In [25]:

```

from scipy import stats

t_value,p_value=stats.ttest_ind(PreLockdownSec,PostLockdownSec)

print('Test statistic is %f'%float("{:.6f}".format(t_value)))

print('p-value for two tailed test is %f'%p_value)

alpha = 0.05

```

```

Test statistic is -13.814048
p-value for two tailed test is 0.000000

```

The p value is less than 0.05 therefore we reject the null hypothesis. Therefore lockdown Black Coffee's spins changed significantly after the lockdown date. This will need to be analysed as its possible that spins did not change due to lockdown, there could have been other factors.

In [26]:

```

# Export data to CSV for use in Tableau
data.to_csv('BlackCoffee_Cleaned_dataset.csv')

```

## Time series analysis of Secs

In [27]:

```

# create a new data set to do a basic TS analysis for Seconds

df = data[['Sec', 'Date']]
df.head()

```

Out[27]:

	Sec	Date
2	180	2020-01-01 00:35:00
5	216	2020-01-01 00:56:00
6	198	2020-01-01 00:57:00
12	198	2020-01-01 04:39:00
15	66	2020-01-01 05:08:00

In [28]:

```
#use the time variable as the new index to the dataframe and drop it from the main a
df.set_index('Date', inplace = True)
df
```

Out[28]:

	Sec
Date	
2020-01-01 00:35:00	180
2020-01-01 00:56:00	216
2020-01-01 00:57:00	198
2020-01-01 04:39:00	198
2020-01-01 05:08:00	66
...	...
2020-12-31 19:44:00	432
2020-12-31 20:16:00	354
2020-12-31 20:31:00	204
2020-12-31 22:09:00	324
2020-12-31 23:08:00	330

8970 rows × 1 columns


In [29]:

```
# Check dataframe info
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 8970 entries, 2020-01-01 00:35:00 to 2020-12-31 23:08:0
0
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  -
0    Sec      8970 non-null     int64
dtypes: int64(1)
memory usage: 140.2 KB
```

In [30]:

```
#plot the series
plt.style.use('ggplot')
plt.rcParams['figure.figsize']=(10,10)
df.plot(linewidth=1, color='black')
plt.show()
```



In [32]:

```
# produce a multiplicative and additive decomposition of the series

sdc_a = seasonal_decompose(dfsum, model='a') # this is the model-parameter's default
sdc_m = seasonal_decompose(dfsum, model='m')

additive_model = sdc_a.trend+sdc_a.seasonal # deterministic part of this decomposition
multiplicative_model = sdc_m.trend*sdc_m.seasonal
```



In [33]:

*#Plot the different seasonal models*

```


plt.rcParams['figure.figsize']=(10,6)
plt.style.use('ggplot')
plt.figure(1)
plt.subplot(2,1,1)
plt.plot(dfsum, label = 'Original', color='grey',linewidth=3)
plt.plot(multiplicative_model, label = 'Multiplicative',linestyle='--',color='blue',
plt.legend(loc = 'best')
plt.subplot(2,1,2)
plt.plot(dfsum, label = 'Original',color='grey')
plt.plot(additive_model, label = 'Additive',linestyle='-.', color='green', linewidth=3)
plt.legend(loc = 'best')
plt.show()

```



In [34]:

```
#plot the two decomposition residuals
plt.style.use('ggplot')
plt.rcParams['figure.figsize']=(14,6)
x = dfsum.index
residual_m = sdc_m.resid
plt.subplot(2,1,1)
plt.title('Multiplicative decomposition residuals')
plt.plot(x, residual_m, 'b', linewidth=2)
residual_a = sdc_a.resid
plt.subplot(2,1,2)
plt.title('Additive decomposition residuals')
plt.plot(x, residual_a, 'r', linewidth=2)
plt.show()
```




In [35]:

```
# use the multiplicative model as variance is smaller
# Take the log of the multiplicative model
df_log = np.log(dfsum)
```

## Acf Plot

In [36]:

```
plt.style.use('fivethirtyeight')
plt.rcParams['figure.figsize']=(15,10)
plot_acf(df_log,lags=40)
plt.show()
```




There is clearly a seasonal component every Saturday. This could mean that stations are playing Black Coffee's tracks for a longer time period on Saturday's and more regularly.

In [37]:

```
# Decompose series into its trend and seasonal parts

res = seasonal_decompose(df_log, model = "multiplicative", period = 30)

fig, (ax1, ax2, ax3) = plt.subplots(3, 1, figsize=(15, 8))
res.trend.plot(ax=ax1, ylabel = "trend")
res.resid.plot(ax=ax2, ylabel = "seasonality")
res.seasonal.plot(ax=ax3, ylabel = "residual")
plt.show()
```



There seems to be a clear upward trend and some seasonality, this can be investigated further if needed.

## Outlier detection

In [112]:

```
# Create a dataframe to check for outliers
df2 = data[['Sec', 'Date']]
df2.rename(columns={"Date": "ds", "Sec": "y"}, inplace=True)
df2.set_index('ds', inplace = True)
#
#df2.resample('D').sum()
#df2['ds'] = pd.to_datetime(df2['ds'])
df2sum = df2.resample('D').sum()
df2sum.reset_index(inplace=True)
df2sum
```

Out[112]:

	ds	y
0	2020-01-01	7818
1	2020-01-02	6156
2	2020-01-03	5958
3	2020-01-04	6714
4	2020-01-05	3294
...	...	...
361	2020-12-27	8190
362	2020-12-28	12228
363	2020-12-29	12042
364	2020-12-30	13410
365	2020-12-31	12978

366 rows × 2 columns



In [119]:

```

#create a plot to show outliers above the 99 percent confidence interval for predict
from fbprophet import Prophet
import altair as alt
def fit_predict_model(dataframe, interval_width = 0.99, changepoint_range = 0.8):
    m = Prophet(daily_seasonality = False, yearly_seasonality = False, weekly_seasonality = False,
                seasonality_mode = 'multiplicative',
                interval_width = interval_width,
                changepoint_range = changepoint_range)
    m = m.fit(dataframe)

    forecast = m.predict(dataframe)
    forecast['fact'] = dataframe['y'].reset_index(drop = True)
    print('Displaying Prophet plot')
    fig1 = m.plot(forecast)
    return forecast
# Display outputs
pred = fit_predict_model(df2sum)

```


Initial log joint probability = -18.6621

	Iter	log prob	dx	grad	alpha	al
pha0	# evals	Notes				
	69	559.264	0.00339809	80.1847	3.42e-05	
0.001	117	LS failed, Hessian reset				
	99	559.598	0.000509851	68.016	0.616	
0.616	154					
	Iter	log prob	dx	grad	alpha	al
pha0	# evals	Notes				
	199	559.676	1.0711e-06	71.3944	0.5954	0.
5954	280					
	Iter	log prob	dx	grad	alpha	al
pha0	# evals	Notes				
	208	559.678	7.81549e-06	72.4875	9.793e-08	
0.001	333	LS failed, Hessian reset				
	293	559.68	2.03671e-09	78.184	0.001314	
1	459					

Optimization terminated normally:

Convergence detected: absolute parameter change was below tolerance

Displaying Prophet plot



In [120]:

```
def detect_anomalies(forecast):
    forecasted = forecast[['ds', 'trend', 'yhat', 'yhat_lower', 'yhat_upper', 'fact']]
    #forecast['fact'] = df['y']

    forecasted['anomaly'] = 0
    forecasted.loc[forecasted['fact'] > forecasted['yhat_upper'], 'anomaly'] = 1
    forecasted.loc[forecasted['fact'] < forecasted['yhat_lower'], 'anomaly'] = -1

    #anomaly importances
    forecasted['importance'] = 0
    forecasted.loc[forecasted['anomaly'] == 1, 'importance'] = \
        (forecasted['fact'] - forecasted['yhat_upper'])/forecasted['fact']
    forecasted.loc[forecasted['anomaly'] == -1, 'importance'] = \
        (forecasted['yhat_lower'] - forecasted['fact'])/forecasted['fact']

    return forecasted

pred = detect_anomalies(pred)
```

In [118]:

```
#create a list of anomalies that lie outside the 99 percent confidence interval
anoms = pred[pred.anomaly == 1]
anoms
```

Out[118]:

	ds	trend	yhat	yhat_lower	yhat_upper	fact	anomaly	importance
<b>129</b>	2020-05-09	5062.055836	5062.055836	-150.786232	9701.056949	9990	1	0.028923
<b>136</b>	2020-05-16	5133.561341	5133.561341	512.659916	10505.086482	10830	1	0.030001
<b>199</b>	2020-07-18	5777.110851	5777.110851	771.365815	11065.441282	12138	1	0.088364
<b>318</b>	2020-11-14	8841.654228	8841.654228	3224.100983	14414.595445	14478	1	0.004379
<b>332</b>	2020-11-28	9389.481523	9389.481523	4085.983047	14380.998292	15090	1	0.046985
<b>339</b>	2020-12-05	9663.395170	9663.395170	4436.628083	14913.172874	15978	1	0.066643
<b>346</b>	2020-12-12	9937.308817	9937.308817	5282.605199	15082.600278	15234	1	0.009938

In [155]:

```
#Check for which day of the week the outliers sit: 0 = Monday, 1 = Tues etc
s = anoms.ds
s.dt.dayofweek
```

Out[155]:

```
129    5
136    5
199    5
318    5
332    5
339    5
346    5
Name: ds, dtype: int64
```

All outliers occur on a Friday, this could be because of there is more and longer play of Black Coffee's music on the start of the weekend, or there could be more party related radio shows on that day